

Code :-

```
# Assignment 4: Demonstrate Python Libraries

import math, numpy as np, pandas as pd

from scipy import stats

import matplotlib.pyplot as plt


# [1] Math Library

print("MATH LIBRARY:")

print("Square root of 25:", math.sqrt(25))

print("Factorial of 5:", math.factorial(5))

print("Ceil(3.2):", math.ceil(3.2), " Floor(3.8):", math.floor(3.8),
"\n")


# [2] NumPy & SciPy

arr = np.random.randint(1, 100, 10)

print("NUMPY ARRAY:", arr)

print("Mean:", np.mean(arr), " Median:", np.median(arr))

print("Std Dev:", np.std(arr), " Var:", np.var(arr))

print("Mode using SciPy:", stats.mode(arr, keepdims=True).mode[0],
"\n")


# [3] Pandas

data = {"Name": ["Aashish", "Ronit", "Vedansh", "Soham"],
        "Age": [20, 21, 22, 23],
        "Marks": [85, 90, 88, 95]}

df = pd.DataFrame(data)

print("PANDAS DATAFRAME:\n", df)

print("\nSummary:\n", df.describe(), "\n")


# [4] Matplotlib Visualization
```

```
plt.figure(figsize=(5,4))

plt.bar(df["Name"] , df["Marks"] , color='orange')

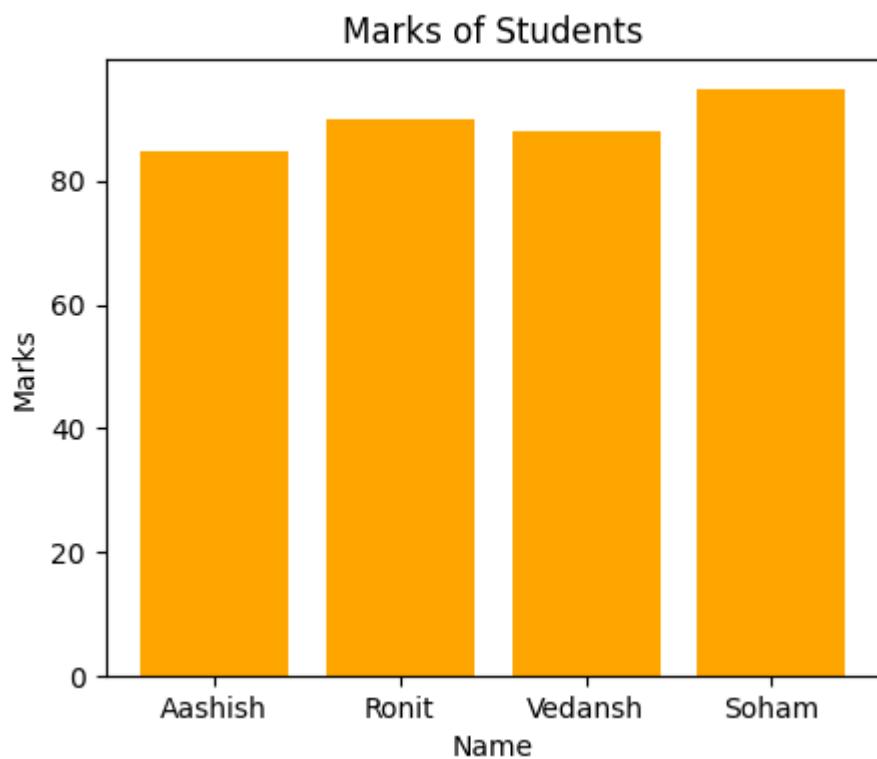
plt.title("Marks of Students")

plt.xlabel("Name")

plt.ylabel("Marks")

plt.show()
```

Output:-



Code:-

```
# assignment5_stats.py
# Basic arithmetic and statistical calculations (Mean, Median, Std Dev)

import numpy as np

# Arithmetic Operations
a = 15
b = 4
print("Arithmetic Operations:")
print(f"Addition: {a} + {b} = {a + b}")
print(f"Subtraction: {a} - {b} = {a - b}")
print(f"Multiplication: {a} * {b} = {a * b}")
print(f"Division: {a} / {b} = {a / b}")
print(f"Modulus: {a} % {b} = {a % b}")
print(f"Exponentiation: {a} ** {b} = {a ** b}")
print()

# Statistical Calculations
data = [10, 20, 30, 40, 50]
mean = np.mean(data)
median = np.median(data)
std_dev = np.std(data)

print("Statistical Calculations:")
print(f"Data: {data}")
print(f"Mean: {mean}")
print(f"Median: {median}")
print(f"Standard Deviation: {std_dev}")
```

Output

```
● (.venv) PS C:\Users\Lenovo\titanic_project> pip install numpy
>>
Requirement already satisfied: numpy in c:\users\lenovo\titanic_project\.venv\lib\site-p
ackages (2.3.4)
● (.venv) PS C:\Users\Lenovo\titanic_project> python assignment5_stats.py
>>
Arithmetic Operations:
Addition: 15 + 4 = 19
Subtraction: 15 - 4 = 11
Multiplication: 15 * 4 = 60
Division: 15 / 4 = 3.75
Modulus: 15 % 4 = 3
Exponentiation: 15 ** 4 = 50625

Statistical Calculations:
Data: [10, 20, 30, 40, 50]
Mean: 30.0
Median: 30.0
Standard Deviation: 14.142135623730951
○ (.venv) PS C:\Users\Lenovo\titanic_project> []
```

Code :-

```
# Assignment 6: Data Preprocessing - Missing Values, Encoding,
Normalization

import pandas as pd
import numpy as np
from sklearn.preprocessing import LabelEncoder, MinMaxScaler,
StandardScaler

# Step 1: Load Dataset (Titanic dataset from CSV or create sample)
try:
    df = pd.read_csv("titanic.csv")      # if you have a Titanic CSV
    print("Loaded titanic.csv successfully!\n")
except:
    print("No titanic.csv found, using sample dataset...\n")
    df = pd.DataFrame({
        "Name": ["Aashish", "Saloni", "Vedansh", "Poornima", "Aditya"],
        "Age": [20, np.nan, 22, 19, 23],
        "Gender": ["Male", "Female", "Male", "Female", "Male"],
        "Salary": [25000, 27000, np.nan, 30000, 24000]
    })

# Step 2: Explore Data
print("Original Data:\n", df, "\n")
print("Missing values per column:\n", df.isnull().sum(), "\n")

# Step 3: Handle Missing Values
df["Age"].fillna(df["Age"].mean(), inplace=True)
df["Salary"].fillna(df["Salary"].mean(), inplace=True)
print("After handling missing values:\n", df, "\n")

# Step 4: Encode Categorical Data
le = LabelEncoder()
df["Gender_Encoded"] = le.fit_transform(df["Gender"])  # Male=1,
                                                       Female=0
print("After Encoding Categorical Data:\n", df, "\n")

# Step 5: Normalization
scaler = MinMaxScaler()
df[["Age_Norm", "Salary_Norm"]] = scaler.fit_transform(df[["Age",
"Salary"]])
print("After Normalization (Min-Max Scaling):\n", df, "\n")

# Step 6: Standardization (optional)
```

```

std_scaler = StandardScaler()
df[["Age_Std", "Salary_Std"]] = std_scaler.fit_transform(df[["Age",
"Salary"]])
print("After Standardization (Z-score):\n", df, "\n")

# Final clean data
print("==== Final Preprocessed Data ====")
print(df)

```

Output :-

The screenshot shows a Jupyter Notebook interface with the 'TERMINAL' tab selected. The terminal output is as follows:

```

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL ...
> < TERMINAL
PS C:\Users\Lenovo\titanic_project> & C:/Users/lenovo/titanic_project/.venv/Scripts/activate.ps1
PS C:\Users\Lenovo\titanic_project> & C:/Users/lenovo/titanic_project/.venv/Scripts/python.exe c:/Users/lenovo/titanic_project/assignment6_preprocessing.py
No titanic.csv found, using sample dataset...

Original Data:
   Name  Age  Gender  Salary
0  Aashish  20.0    Male  25000.0
1  Saloni     NaN  Female  27000.0
2  Vedansh  22.0    Male      NaN
3  Poornima  19.0  Female  30000.0
4   Aditya  23.0    Male  24000.0

Missing values per column:
Name      0
Age       1
Gender     0
Salary     1
dtype: int64

c:\Users\Lenovo\titanic_project\assignment6_preprocessing.py:24: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```

The screenshot shows a Jupyter Notebook interface with the 'TERMINAL' tab selected. The terminal output is as follows:

```

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL ...
> < TERMINAL
Python + ... Python ...
PS C:\Users\Lenovo\titanic_project> & C:/Users/lenovo/titanic_project/.venv/Scripts/python.exe c:/Users/lenovo/titanic_project/assignment6_preprocessing.py

df["Age"].fillna(df["Age"].mean(), inplace=True)
c:\Users\lenovo\titanic_project\assignment6_preprocessing.py:25: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```

```
df["Salary"].fillna(df["Salary"].mean(), inplace=True)
After handling missing values:
      Name   Age  Gender  Salary
0  Aashish  20.0    Male  25000.0
1  Saloni  21.0  Female  27000.0
2  Vedansh  22.0    Male  26500.0
3 Poornima  19.0  Female  30000.0
4  Aditya  23.0    Male  24000.0

After Encoding Categorical Data:
      Name   Age  Gender  Salary  Gender_Encoded
0  Aashish  20.0    Male  25000.0                  1
1  Saloni  21.0  Female  27000.0                  0
2  Vedansh  22.0    Male  26500.0                  1
3 Poornima  19.0  Female  30000.0                  0
4  Aditya  23.0    Male  24000.0                  1

After Normalization (Min-Max Scaling):
      Name   Age  Gender  Salary  Gender_Encoded  Age_Norm  Salary_Norm
0  Aashish  20.0    Male  25000.0                  1       0.25     0.166667
1  Saloni  21.0  Female  27000.0                  0       0.50     0.500000
2  Vedansh  22.0    Male  26500.0                  1       0.75     0.416667
3 Poornima  19.0  Female  30000.0                  0       0.00     1.000000
4  Aditya  23.0    Male  24000.0                  1       1.00     0.000000
```

```

After Standardization (Z-score):
      Name  Age  Gender  Salary  ...  Age_Norm  Salary_Norm  Age_Std  Salary_Std
0   Aashish  20.0    Male  25000.0  ...     0.25    0.166667 -0.707107 -0.731925
1   Saloni  21.0  Female  27000.0  ...     0.50    0.500000  0.000000  0.243975
2   Vedansh 22.0    Male  26500.0  ...     0.75    0.416667  0.707107  0.000000
3  Poornima  19.0  Female  30000.0  ...     0.00    1.000000 -1.414214  1.707825
4   Aditya  23.0    Male  24000.0  ...     1.00    0.000000  1.414214 -1.219875

[5 rows x 9 columns]
4   Aditya  23.0    Male  24000.0           1

After Normalization (Min-Max Scaling):
      Name  Age  Gender  Salary  Gender_Encoded  Age_Norm  Salary_Norm
0   Aashish  20.0    Male  25000.0            1       0.25    0.166667
1   Saloni  21.0  Female  27000.0            0       0.50    0.500000
2   Vedansh 22.0    Male  26500.0            1       0.75    0.416667
3  Poornima  19.0  Female  30000.0            0       0.00    1.000000
4   Aditya  23.0    Male  24000.0            1       1.00    0.000000

○ After Standardization (Z-score):
      Name  Age  Gender  Salary  ...  Age_Norm  Salary_Norm  Age_Std  Salary_Std
0   Aashish  20.0    Male  25000.0  ...     0.25    0.166667 -0.707107 -0.731925
1   Saloni  21.0  Female  27000.0  ...     0.50    0.500000  0.000000  0.243975
2   Vedansh 22.0    Male  26500.0  ...     0.75    0.416667  0.707107  0.000000
3  Poornima  19.0  Female  30000.0  ...     0.00    1.000000 -1.414214  1.707825

```

```

--- Final Preprocessed Data ---
      Name  Age  Gender  Salary  ...  Age_Norm  Salary_Norm  Age_Std  Salary_Std

--- Final Preprocessed Data ---
      Name  Age  Gender  Salary  ...  Age_Norm  Salary_Norm  Age_Std  Salary_Std
0   Aashish  20.0    Male  25000.0  ...     0.25    0.166667 -0.707107 -0.731925
1   Saloni  21.0  Female  27000.0  ...     0.50    0.500000  0.000000  0.243975
      Name  Age  Gender  Salary  ...  Age_Norm  Salary_Norm  Age_Std  Salary_Std
0   Aashish  20.0    Male  25000.0  ...     0.25    0.166667 -0.707107 -0.731925
1   Saloni  21.0  Female  27000.0  ...     0.50    0.500000  0.000000  0.243975
0   Aashish  20.0    Male  25000.0  ...     0.25    0.166667 -0.707107 -0.731925
1   Saloni  21.0  Female  27000.0  ...     0.50    0.500000  0.000000  0.243975
1   Saloni  21.0  Female  27000.0  ...     0.50    0.500000  0.000000  0.243975
2   Vedansh 22.0    Male  26500.0  ...     0.75    0.416667  0.707107  0.000000
3  Poornima  19.0  Female  30000.0  ...     0.00    1.000000 -1.414214  1.707825
4   Aditya  23.0    Male  24000.0  ...     1.00    0.000000  1.414214 -1.219875

[5 rows x 9 columns]
(.venv) PS C:\Users\Lenovo\titanic_project>

```

Code:

```
# Step 1: Import libraries

import numpy as np
import pandas as pd
from sklearn.datasets import load_breast_cancer
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix, roc_curve, auc
import matplotlib.pyplot as plt
import seaborn as sns

# Step 2: Load the dataset
data = load_breast_cancer()
X = data.data
y = data.target

# Step 3: Split into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Step 4: Train Logistic Regression model
model = LogisticRegression(max_iter=10000)
model.fit(X_train, y_train)

# Step 5: Predictions
y_pred = model.predict(X_test)
y_proba = model.predict_proba(X_test)[:, 1] # Probabilities for ROC curve

# Step 6: Evaluation
```

```
# Accuracy

acc = accuracy_score(y_test, y_pred)
print(f"Accuracy: {acc:.4f}")


# Confusion Matrix

cm = confusion_matrix(y_test, y_pred)
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.title("Confusion Matrix")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()


# ROC Curve

fpr, tpr, thresholds = roc_curve(y_test, y_proba)
roc_auc = auc(fpr, tpr)

plt.figure()
plt.plot(fpr, tpr, color='darkorange', lw=2, label=f'ROC curve (AUC = {roc_auc:.2f})')
plt.plot([0, 1], [0, 1], color='gray', linestyle='--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic')
plt.legend(loc="lower right")
plt.grid(True)
plt.show()
```

Output:

Accuracy: 0.9561

Figure 1

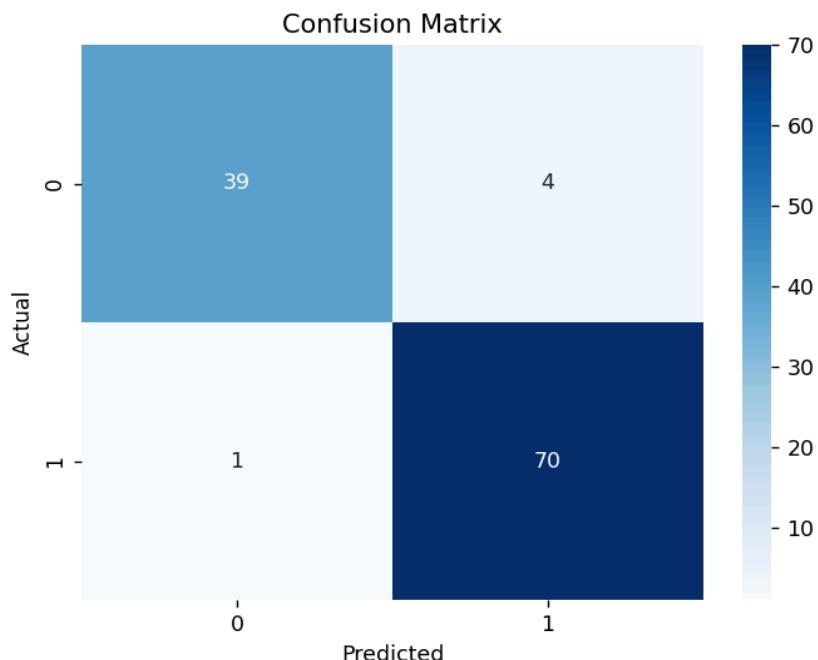
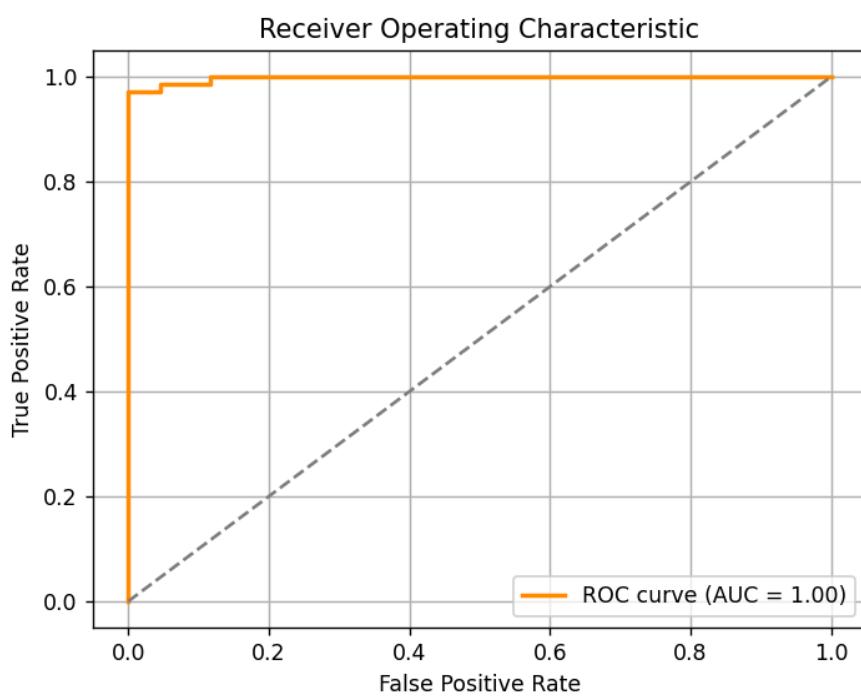


Figure 1



Code:

```
# Step 1: Import libraries
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
import seaborn as sns
import matplotlib.pyplot as plt

# Step 2: Load dataset
iris = load_iris()
X = iris.data
y = iris.target

# Step 3: Split into train and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

# Step 4: Train Naive Bayes model
model = GaussianNB()
model.fit(X_train, y_train)

# Step 5: Make predictions
y_pred = model.predict(X_test)

# Step 6: Evaluate the model
# Accuracy
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy:.4f}")
```

```
# Confusion Matrix

cm = confusion_matrix(y_test, y_pred)

sns.heatmap(cm,      annot=True,      cmap='Blues',      xticklabels=iris.target_names,
yticklabels=iris.target_names)

plt.xlabel("Predicted")

plt.ylabel("Actual")

plt.title("Confusion Matrix")

plt.show()

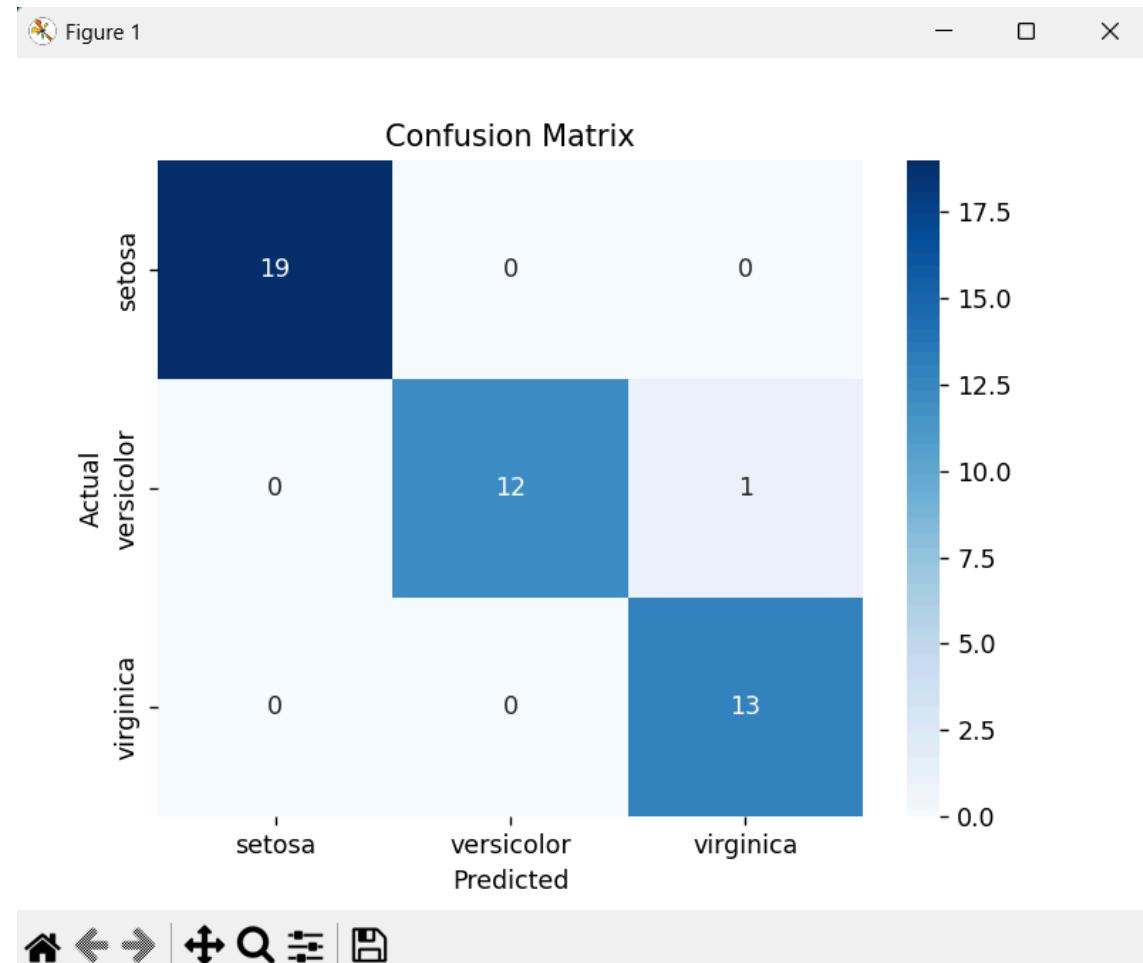
# Precision, Recall, F1-score

print("\n Classification Report:\n")

print(classification_report(y_test, y_pred, target_names=iris.target_names))
```

Output:

Accuracy: 0.9778



Classification Report:

	precision	recall	f1-score	support
setosa	1.00	1.00	1.00	19
versicolor	1.00	0.92	0.96	13
virginica	0.93	1.00	0.96	13
accuracy		0.98	0.98	45
macro avg	0.98	0.97	0.97	45
weighted avg	0.98	0.98	0.98	45

Code:

```
# Step 1: Import required libraries
from sklearn.datasets import load_wine
from sklearn.preprocessing import StandardScaler
from sklearn.cluster import KMeans
from sklearn.decomposition import PCA
from sklearn.metrics import silhouette_score
import matplotlib.pyplot as plt
import seaborn as sns

# Step 2: Load the dataset
data = load_wine()
X = data.data
feature_names = data.feature_names

# Step 3: Normalize the data
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Step 4: Apply K-Means clustering
k = 3 # since wine dataset has 3 classes
kmeans = KMeans(n_clusters=k, random_state=42)
labels = kmeans.fit_predict(X_scaled)

# Step 5: Evaluate with Silhouette Score
sil_score = silhouette_score(X_scaled, labels)
print(f"Silhouette Score: {sil_score:.4f}")
```

```
# Step 6: Visualize clusters using PCA

pca = PCA(n_components=2)

X_pca = pca.fit_transform(X_scaled)

plt.figure(figsize=(8, 6))

sns.scatterplot(x=X_pca[:, 0], y=X_pca[:, 1], hue=labels, palette='Set1')

plt.title("K-Means Clustering (Wine Dataset) - PCA Projection")

plt.xlabel("PCA Component 1")

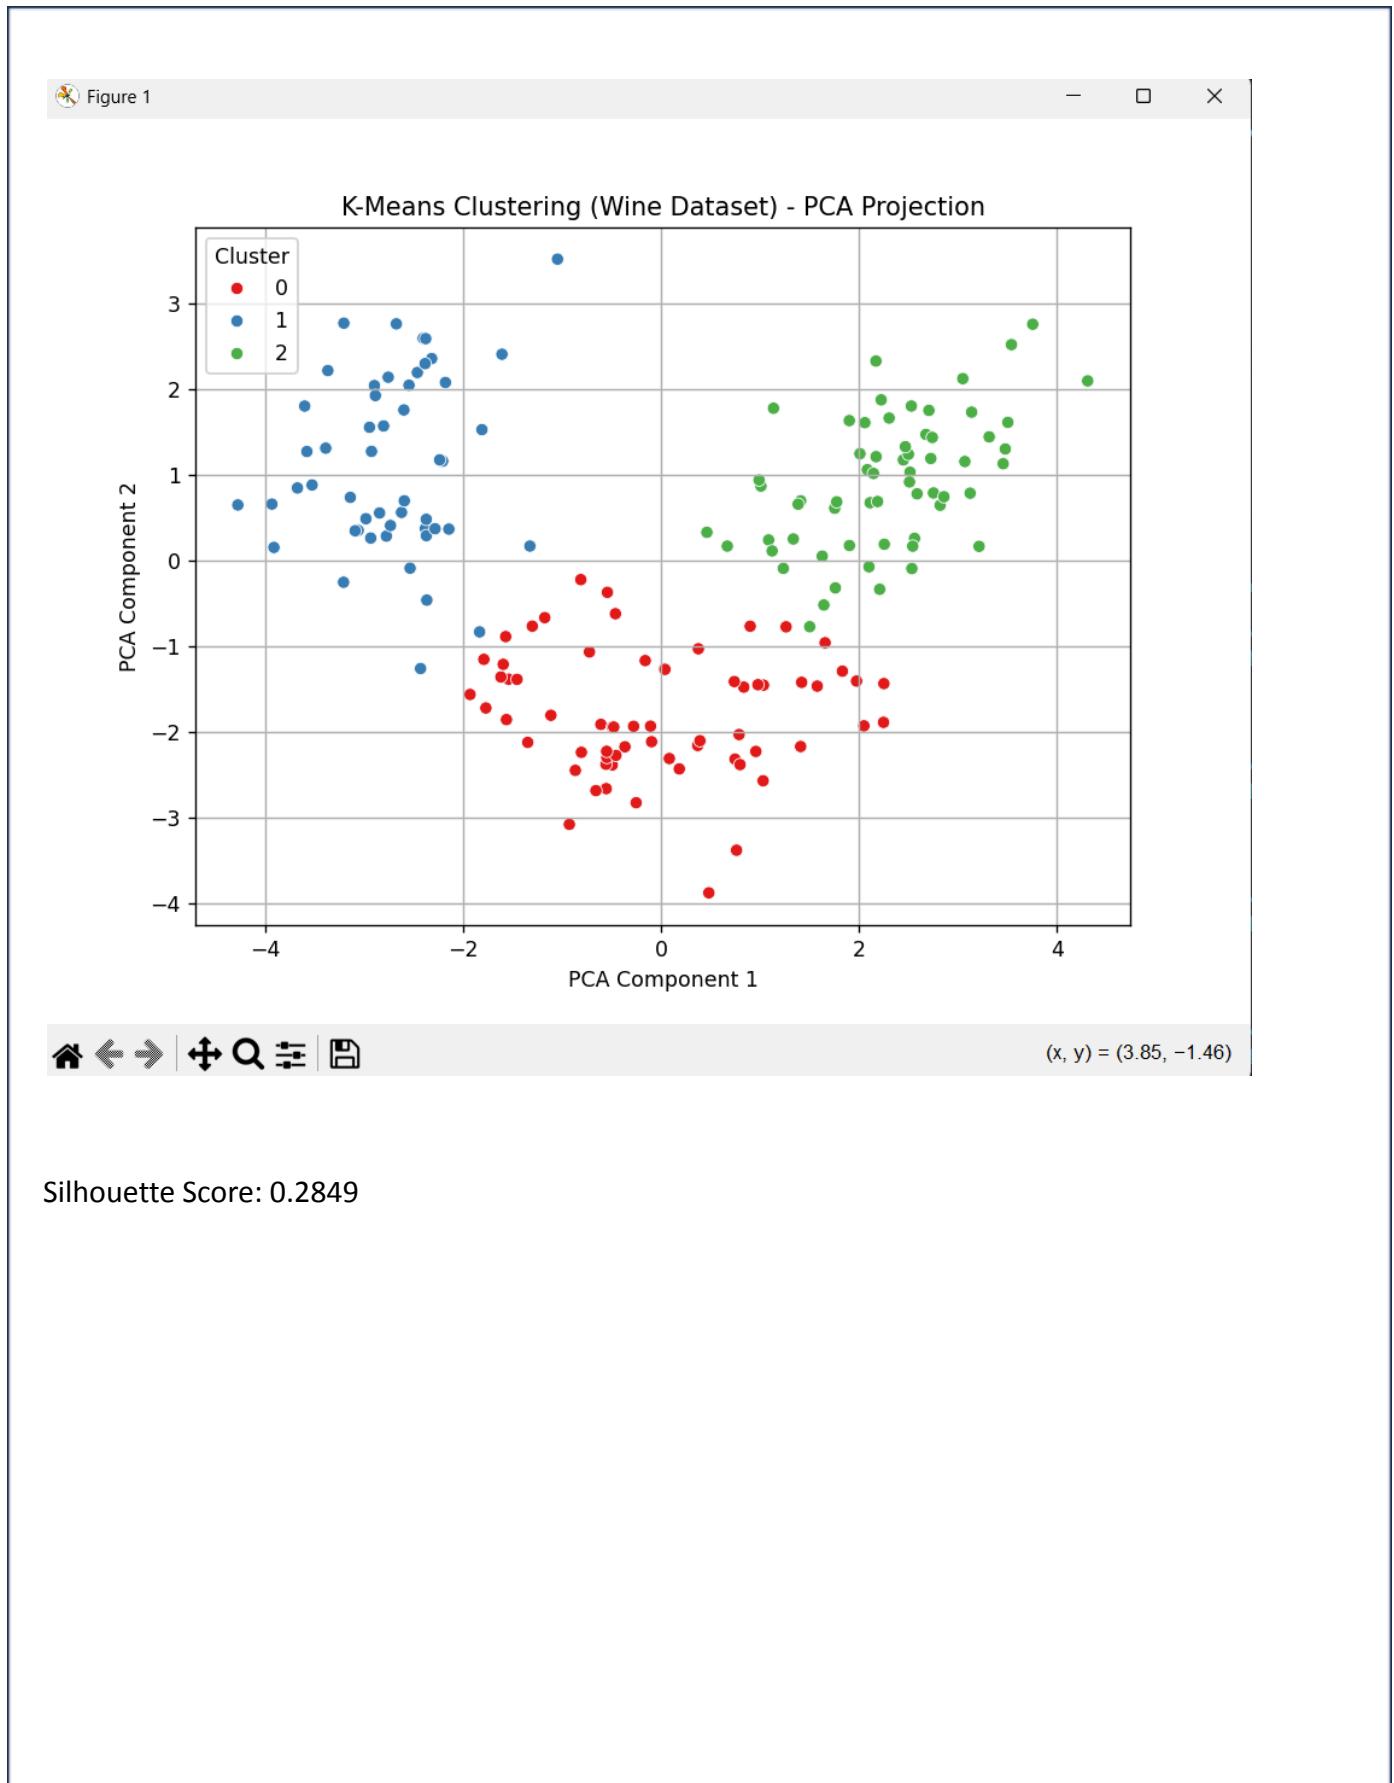
plt.ylabel("PCA Component 2")

plt.grid(True)

plt.legend(title='Cluster')

plt.show()
```

Output:



Code:

```
# Step 1: Import libraries
from sklearn.datasets import load_wine
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
from sklearn.cluster import KMeans
from sklearn.metrics import silhouette_score
import matplotlib.pyplot as plt
import seaborn as sns

# Step 2: Load Wine dataset
data = load_wine()
X = data.data
y = data.target # true labels (for comparison only)

# Step 3: Standardize the data
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Step 4: Apply PCA to reduce to 2 dimensions
pca = PCA(n_components=2)
X_pca = pca.fit_transform(X_scaled)

print("Explained Variance Ratio:", pca.explained_variance_ratio_)
print("Total Variance Retained:", sum(pca.explained_variance_ratio_))

# Step 5: Apply K-Means on 2D PCA data
kmeans = KMeans(n_clusters=3, random_state=42)
```

```
cluster_labels = kmeans.fit_predict(X_pca)

# Step 6: Visualize Clusters
plt.figure(figsize=(8, 6))
sns.scatterplot(x=X_pca[:, 0], y=X_pca[:, 1], hue=cluster_labels, palette="Set1", style=y)
plt.title("K-Means Clustering on PCA-Reduced Wine Dataset")
plt.xlabel("Principal Component 1")
plt.ylabel("Principal Component 2")
plt.grid(True)
plt.legend(title='Cluster / True Class')
plt.show()

# Step 7: Evaluate cluster quality
score = silhouette_score(X_pca, cluster_labels)
print(f"Silhouette Score: {score:.4f}")
```

Output:

Explained Variance Ratio: [0.36198848 0.1920749]

Total Variance Retained: 0.5540633835693526

